

Reader prefetch (#700 Transaction API) on the compressed GPU vector — A100

Prefetch-ahead wins once a slow tier is modeled, peaking where $\text{io} \approx \text{compute}$; at $\text{slow}=0$ the reader is compute-bound (compression already removed the I/O).

