


Eternia vs GPU-memory baselines on ONE 8 GiB GPU (real SIFT k-means)



Only Eternia and UVM run past HBM on one GPU. Eternia is ~60x faster than UVM here (3.4x less data moved + controlled streaming vs page-fault paging). UVM/NVSHMEM/MPI are WSL2-distorted or IPC-blocked — trustworthy baseline numbers need a native-Linux GPU (Delta).