

NeuroPress v2 (model B): predicted vs measured, one point per (chunk, setting)

Top: held-out predictions from cross-validation on the synthetic training corpus (15,000-point sample shown; errors over all 441,000).

Bottom: the saved model on real VPIC fields, never trained on.

