

NeuroPress v2 inside Clio: per-chunk prediction error vs what Clio measured, as the stream goes on

Each point is one chunk (the setting v2 picked, predicted before learning from that chunk), against what Clio measured.
Lines: MAPE over a rolling window of 25 chunks; the number at the right is the MAPE at the end of the stream. Learning rate 0.5; exploration K = 3.

